


# Virtual Memory



- Extending the mem hierarchy
- Virtual Memory
- MMU; Page map
- Page faults
- TLB

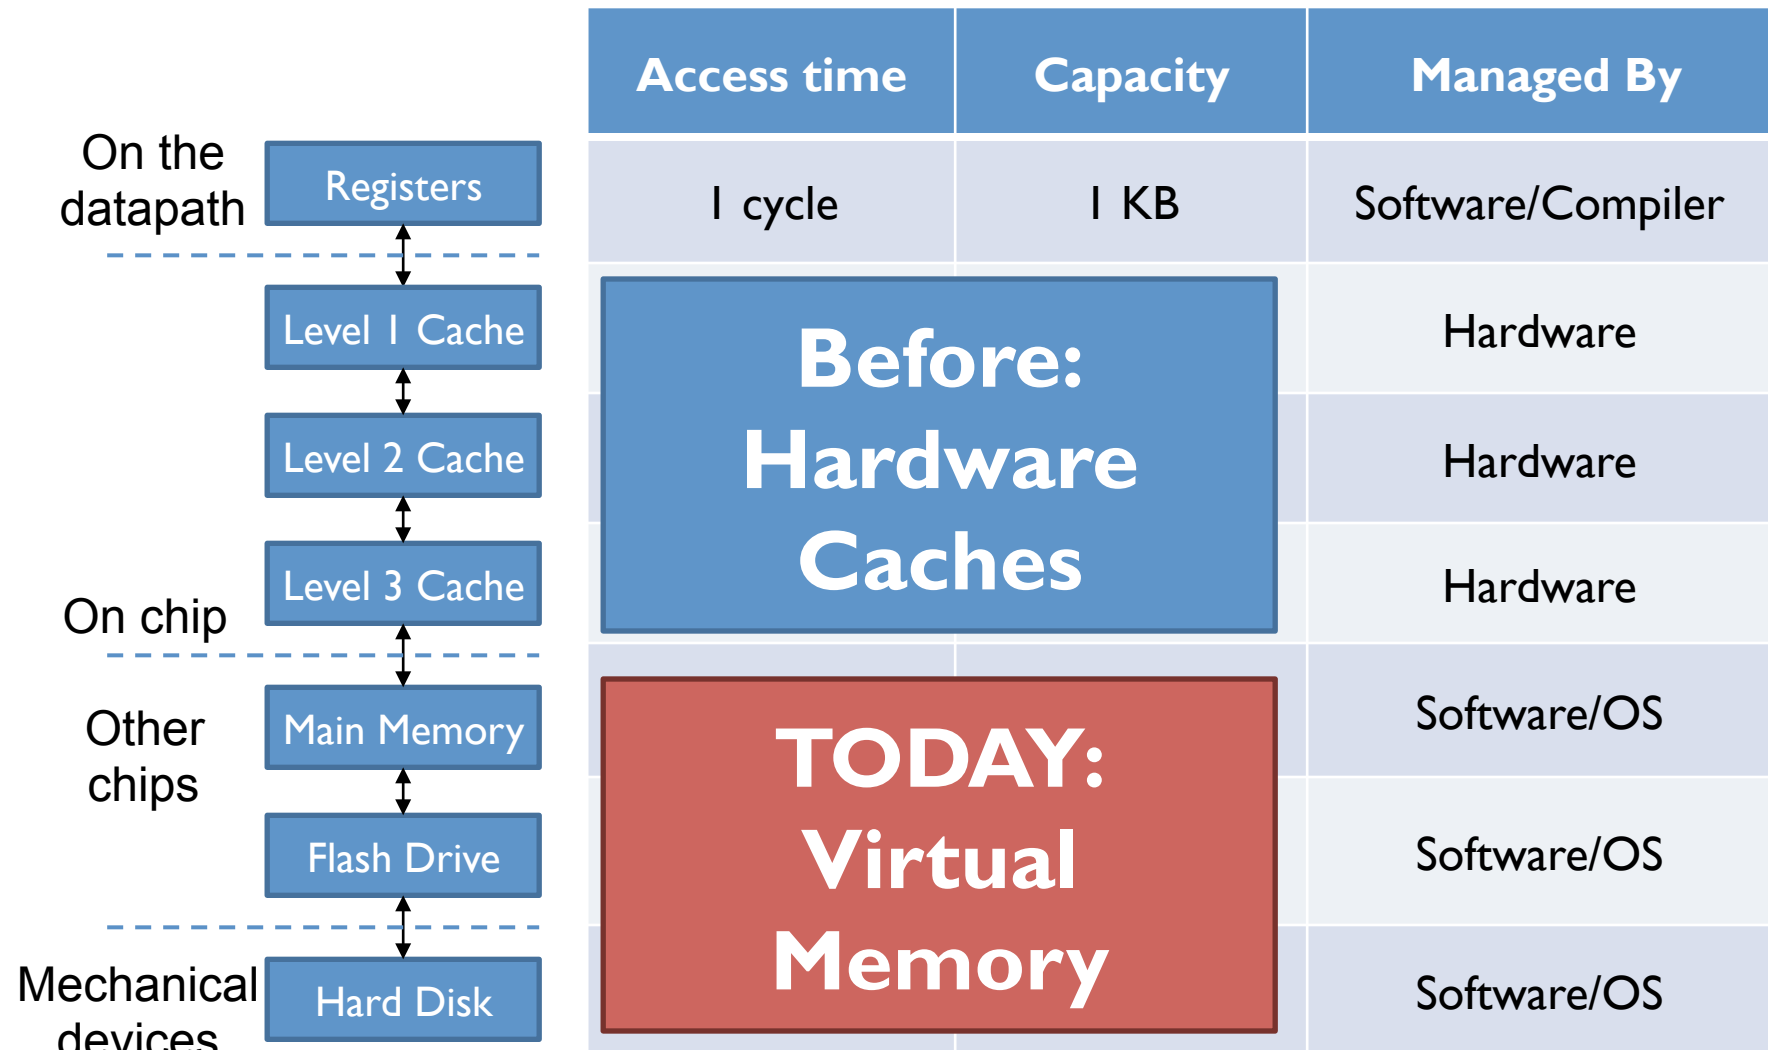
Today's handouts:

- Lecture slides

**Quiz 3 this Friday**


# Reminder: A Typical Memory Hierarchy

- Everything is a cache for something else



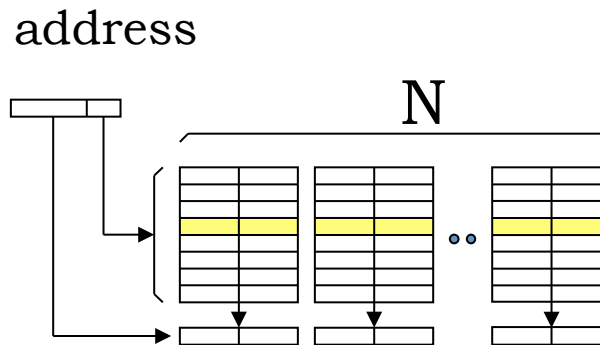
# Reminder: Hardware Caches

Direct-mapped




Each address maps to a single location in the cache, given by index bits

N-way set-associative



Each address can map to any of N locations (ways) of a single set, given by its index bits

Fully associative




Any location can map to any address

## Other implementation choices:

- Block size
- Replacement policy (LRU, Random, ...)
- Write policy (write-through, write-behind, write-back)

# Extending the Memory Hierarchy



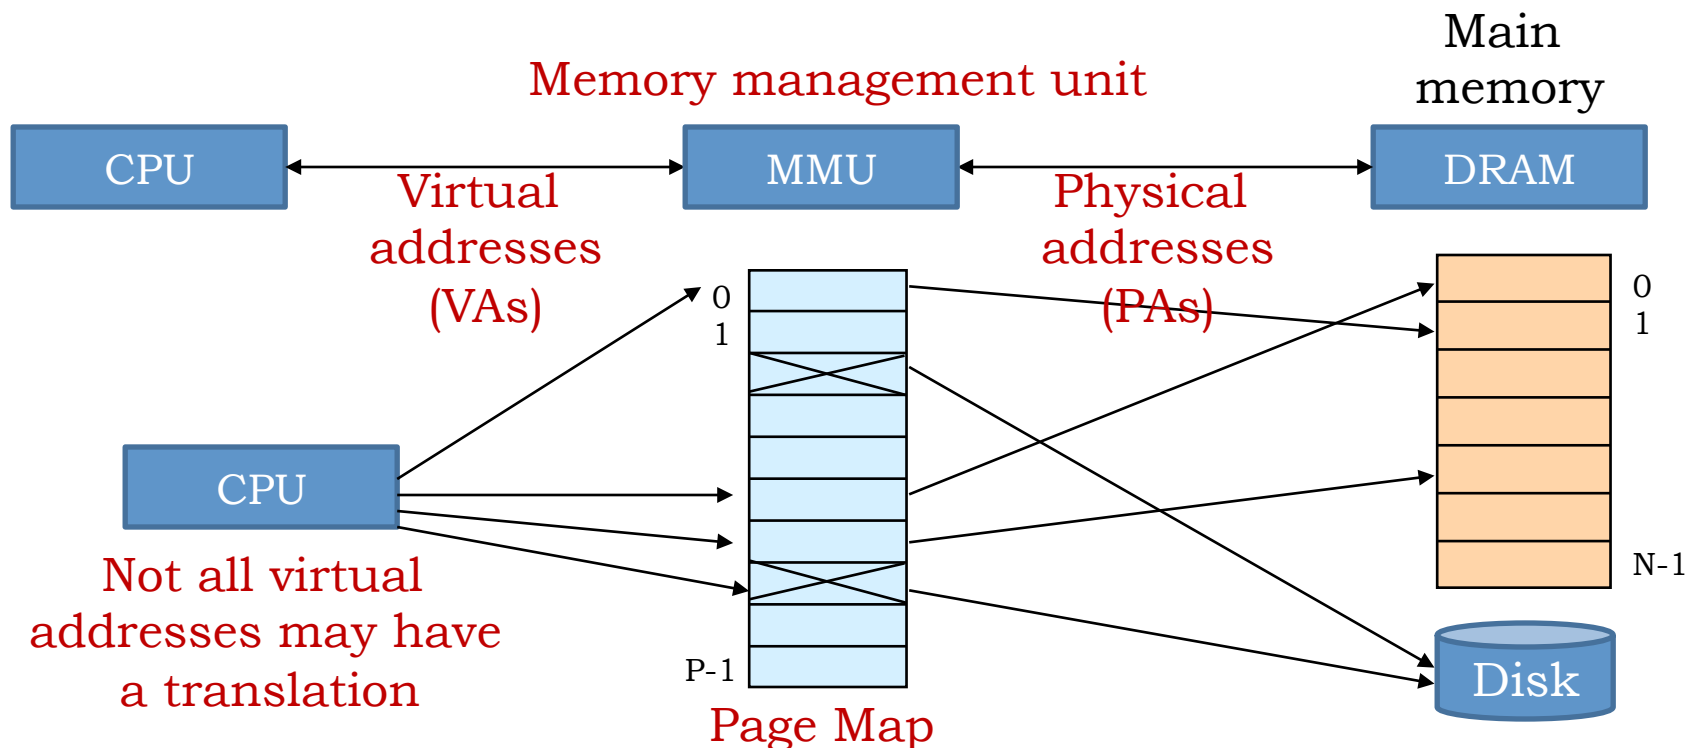
- Problem: DRAM vs disk has much more extreme differences than SRAM vs DRAM
  - Access latencies:
    - DRAM ~10-100x slower than SRAM
    - Disk ~100,000x slower than DRAM
  - Importance of sequential accesses:
    - DRAM: Fetching successive words ~5x faster than first word
    - Disk: Fetching successive words ~100,000x faster than first word
- Result: Design decisions driven by enormous cost of misses

# Impact of Enormous Miss Penalty

- If DRAM was to be organized like an SRAM cache, how should we design it?
  - Associativity: High, minimize miss ratio
  - Block size: Large, amortize cost of a miss over multiple words (locality)
  - Write policy: Write back, minimize number of writes
- Is there anything good about misses being so expensive?
  - We can handle them in software! What's 1000 extra instructions ( $\sim 1\mu\text{s}$ ) vs 10ms?
  - Approach: Handle hits in hardware, misses in software
    - Simpler implementation, more flexibility

# Virtual Memory

- Two kinds of addresses:
  - CPU uses **virtual addresses**
  - Main memory uses **physical addresses**
- Hardware translates virtual addresses to physical addresses via an operating system (OS)-managed table, the **page map**




# Virtual Memory Goals

- Why an extra level of indirection?
- Why let OS control the virtual-physical address mapping?
- Virtual memory goals:
  - Using main memory as disk cache
  - Protecting multiple programs from each other
  - Easing memory management

# Virtual Memory Implementation: Paging

- Divide physical memory in fixed-size blocks, called **pages**
  - Typical page size ( $2^p$ ): 4-16 KB
  - Virtual address: Virtual page number + offset bits
  - Physical address: Physical page number + offset bits
  - Why use lower bits as offset?
- MMU maps virtual address to physical address, or causes a **page fault** (a miss) if no translation
  - Use page map to perform translation




## Terminology

Caching disk using  
main memory =  
*paging or demand  
paging*




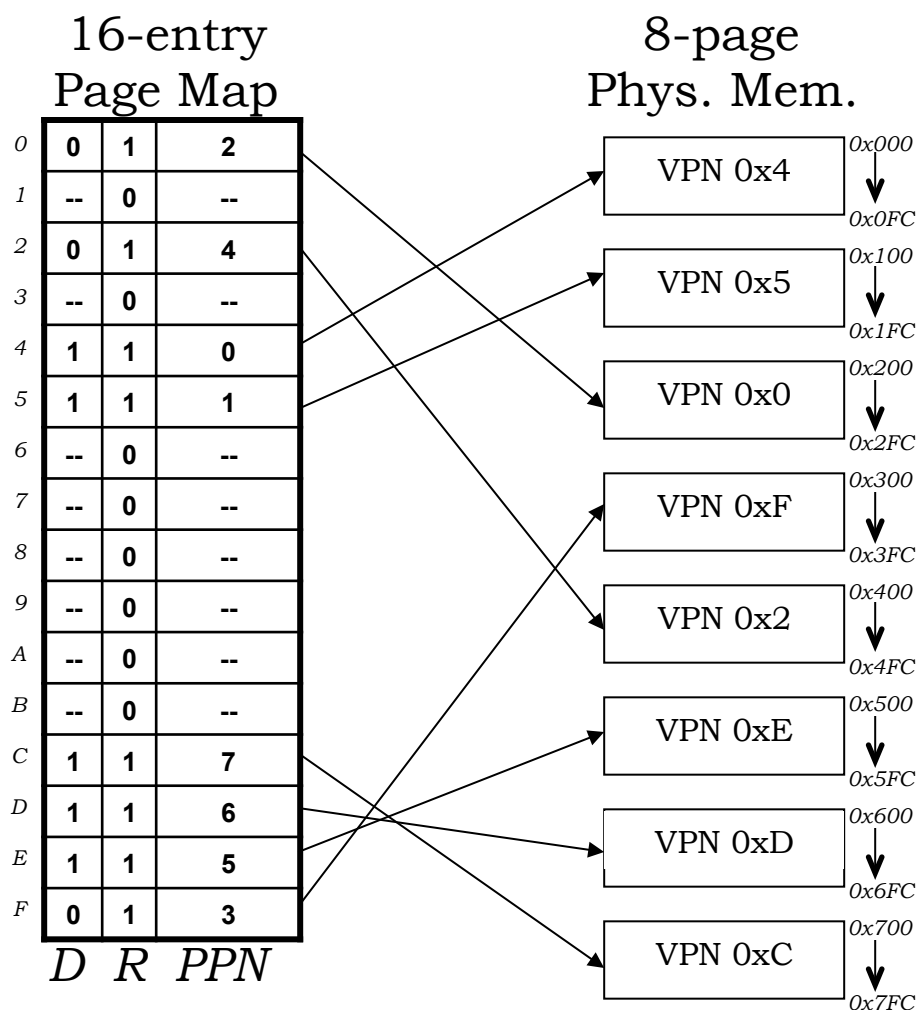
# Simple Page Map Design



## One entry per virtual page

- Resident bit  $R = 1$  for pages stored in RAM, or 0 for non-resident (disk or unallocated). Page fault when  $R = 0$
- Contains physical page number (PPN) of each resident page
- Dirty bit says if we've changed this page since loading it from disk (and therefore need to write it to disk when it's replaced)

# Example: Virtual → Physical Translation



Setup:

256 bytes/page ( $2^8$ )

16 virtual pages ( $2^4$ )

8 physical pages ( $2^3$ )

12-bit VA (4 vpn, 8 offset)

11-bit PA (3 ppn, 8 offset)

LRU page: VPN = 0xE

LD(R31,0x2C8,R0):

VA = 0x2C8, PA = 0x4C8


VPN = 0x2

→ PPN = 0x4


# Page Faults

- If a page does not have a valid translation, MMU causes a **page fault**, a specific type of **exception**
- OS exception handler is invoked, handles miss:
  - Fetch page from disk
  - Choose a page to replace, write it back if dirty
    - Are there any restrictions on which page we can we select?
  - Place new page in memory and update page map
  - Return control to program, which re-executes memory access

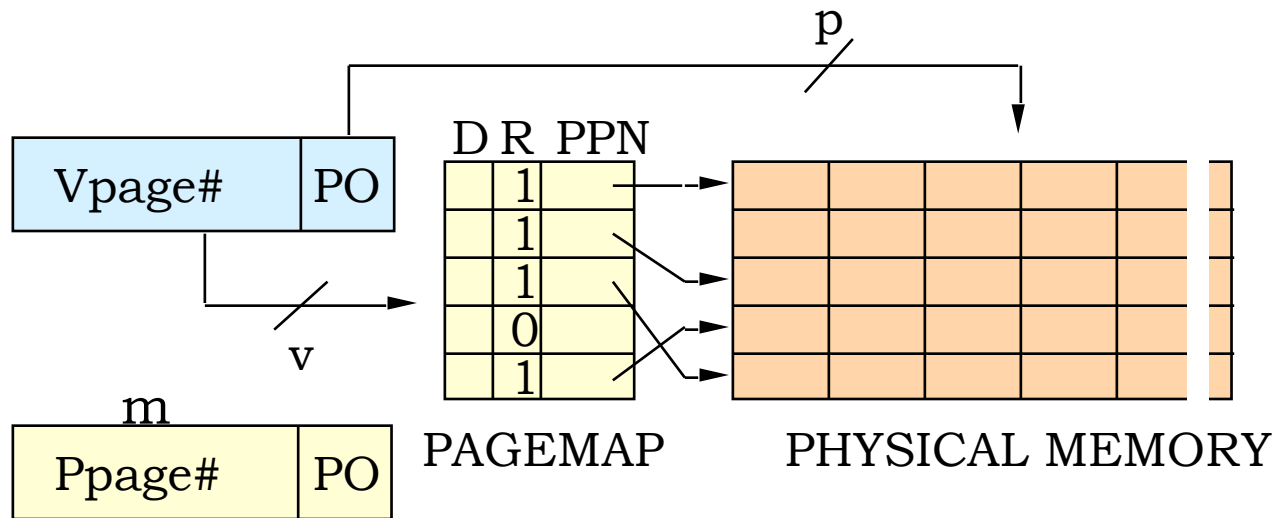
## Before Page Fault



## After Page Fault



# Page Map Arithmetic




$(v + p)$  bits in virtual address  
 $(m + p)$  bits in physical address  
 $2^v$  number of VIRTUAL pages  
 $2^m$  number of PHYSICAL pages  
 $2^p$  bytes per physical page  
 $2^{v+p}$  bytes in virtual memory  
 $2^{m+p}$  bytes in physical memory  
 $(m+2)2^v$  bits in the page map

Typical page size: 4-16 KB  
 Typical  $(v+p)$ : 32-64 bits  
 (4GB-16EB)  
 Typical  $(m+p)$ : 30-40 bits  
 (1GB-1TB)

Long virtual addresses allow ISAs to support larger memories → ISA longevity

# Example: Page Map Arithmetic



SUPPOSE...

32-bit Virtual address

$2^{12}$  page size (4 KB)

$2^{30}$  RAM max (1 GB)

THEN:

# Physical Pages =  $2^{18} = 256K$

# Virtual Pages =  $2^{20}$

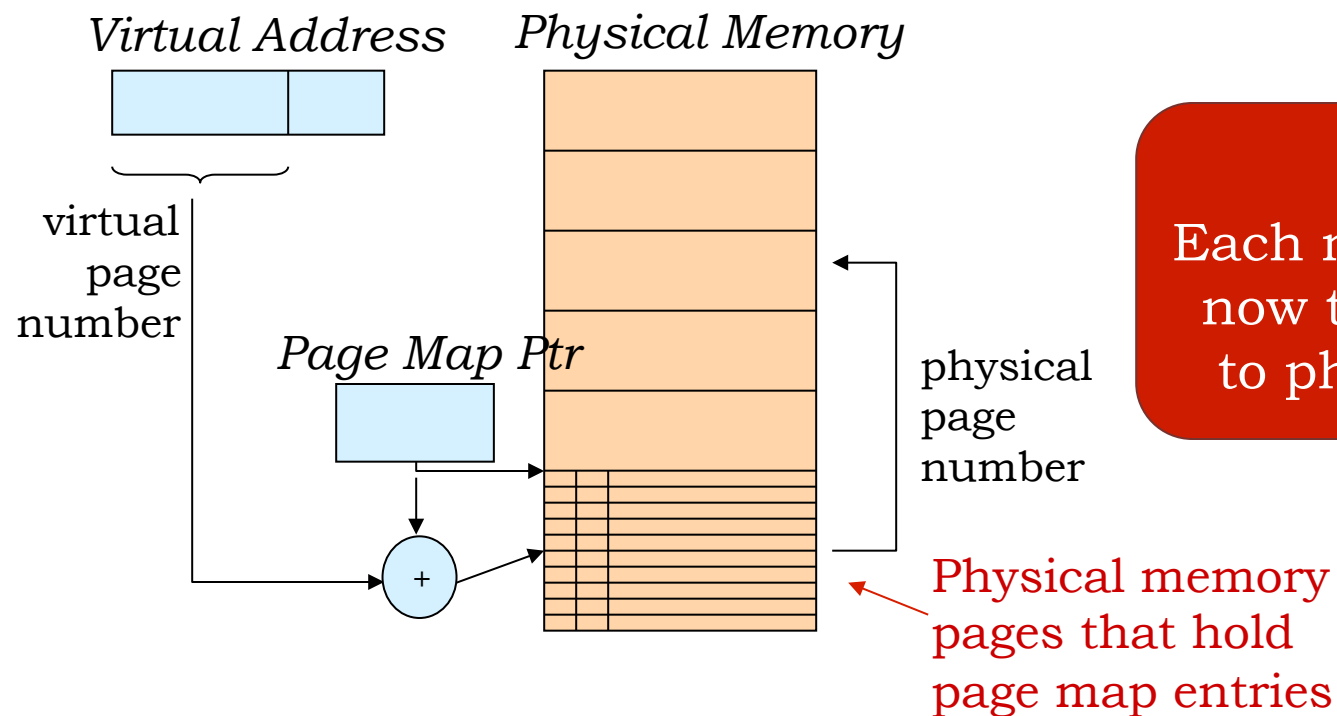
# Page Map Entries =  $2^{20} = 1M$

# Bits In pagemap =  $20 * 2^{20} \sim 20M$

Use fast SRAM for page map??? **OUCH!**

# RAM-Resident Page Maps


- **Small** page maps can use dedicated SRAM... gets expensive for big ones!
- Solution: Move page map to **main memory**:



**PROBLEM**  
Each memory reference  
now takes 2 accesses  
to physical memory!

# Translation Look-aside Buffer (TLB)

- Problem: 2x performance hit... each memory reference now takes 2 accesses!
- Solution: **Cache the page map entries**



IDEA:

LOCALITY in memory  
reference patterns →  
SUPER locality in  
references to page  
map

VARIATIONS:

- multi-level page map
- paging the page map!

**TLB: small cache of page table entries**  
**Associative lookup by VPN**

# Putting it All Together: MMU with TLB

Suppose

- virtual memory of  $2^{32}$  bytes
- physical memory of  $2^{24}$  bytes
- page size is  $2^{10}$  (1 K) bytes
- 4-entry fully associative TLB

TLB				
Tag		Data		
VPN		R	D	PPN
0		0	0	3
6		1	1	2
1		1	1	9
3		0	0	5


Page Map				
VPN	R	D	PPN	
0	0	0	7	
1	1	1	9	
2	1	0	0	
3	0	0	5	
4	1	0	5	
5	0	0	3	
6	1	1	2	
7	1	0	4	
8	1	0	1	
			...	

1. How many pages can be stored in physical memory at once?  $2^{14}$
2. How many entries are there in the page table?  $2^{22}$
3. How many bits per entry in the page table? (Assume each entry has PPN, resident bit, dirty bit) 16
4. How many pages does the page table require?  $2^{23}$  bytes =  $2^{13}$  pages
5. What fraction of virtual memory that can be resident?  $1/2^8$
6. What is the physical address for virtual address 0x1804? What components are involved in the translation? [VPN=6] 0x804
7. Same for 0x1080 [VPN=4] 0x1480
8. Same for 0x0FC [VPN=0] page fault




# Contexts

A context is a mapping of VIRTUAL to PHYSICAL locations, as dictated by contents of the page map:




Several programs may be simultaneously loaded into main memory, each in its separate context:

“Context switch”:  
reload the page map?



# Contexts: A Sneak Preview




First Glimpse of a  
VIRTUAL MACHINE

1. TIMESHARING among several programs
  - Separate context for each program
  - OS loads appropriate context into page map when switching among programs
2. Separate context for OS “Kernel” (e.g., interrupt handlers)...
  - “Kernel” vs “User” contexts
  - Switch to Kernel context on interrupt;
  - Switch back on interrupt return.

# Memory Management & Protection

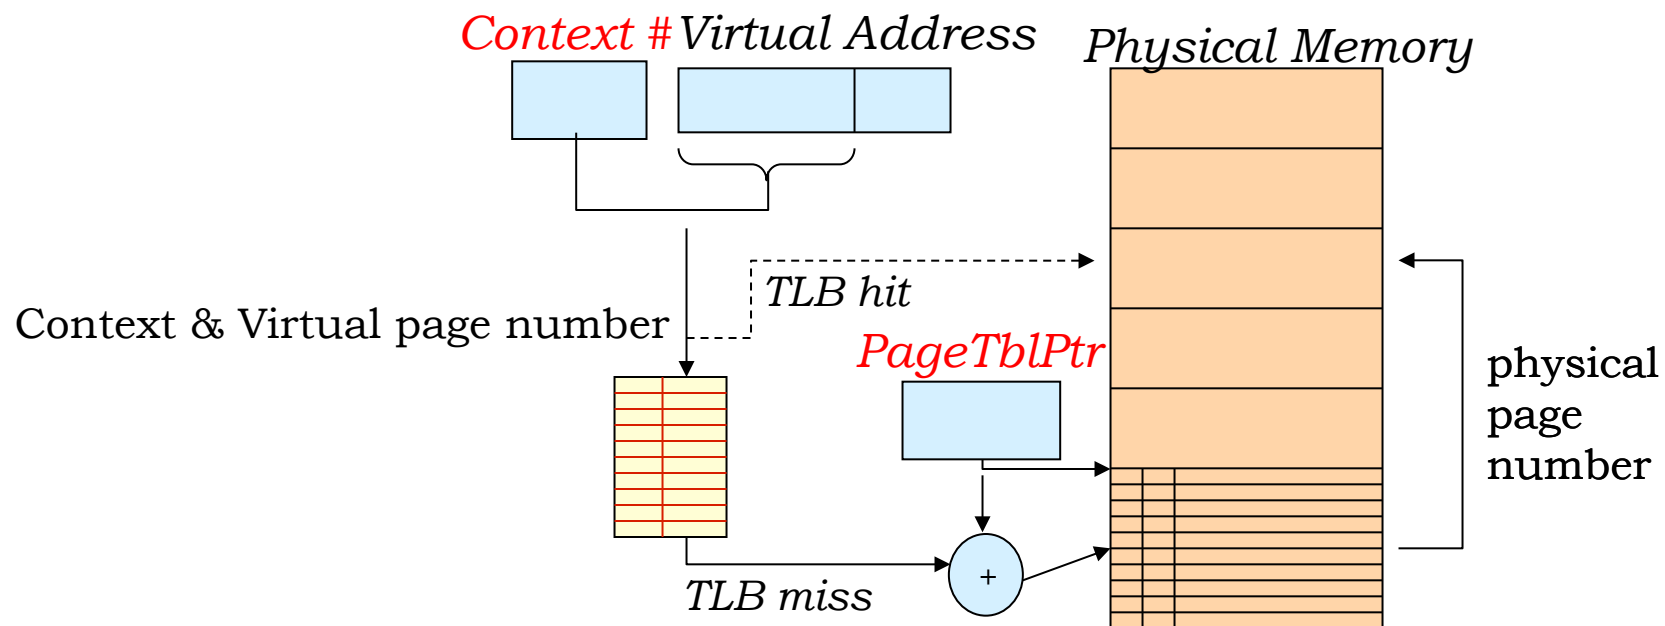
- Applications can be written as if they have access to all memory, without considering where other applications reside
  - Enables fixed conventions (e.g., program starts at 0x1000, stack is contiguous and grows up, ...) without worrying about conflicts
- OS Kernel controls all contexts, prevents programs from reading and writing into each other's memory

*Address Space*



# Rapid Context-Switching


Add a register to hold index of current context. To switch contexts: update Context # and PageTblPtr registers. Don't have to flush TLB since each entry's tag includes context # in addition to virtual page number



# Using Caches with Virtual Memory

## *Virtually-Addressed* Cache


Tags from virtual addresses



- FAST: No MMU time on HIT
- Problem: Must flush cache after context switch


## *Physically-Addressed* Cache

Tags from physical addresses



- Avoids stale cache data after context switch
- SLOW: MMU time on HIT

# Best of Both Worlds: Physically-Indexed, Virtually-Tagged Cache



OBSERVATION: If cache index bits are a subset of page offset bits, tag access in a physical cache can *overlap* page map access. Tag from cache is compared with physical page address from MMU to determine hit/miss.

Problem: Limits #bits of cache index  $\rightarrow$  more size needs higher associativity

# Summary: Virtual Memory

- Goal 1: Exploit locality on a large scale
  - Programmers want a large, flat address space, but use a small portion!
  - Solution: Cache working set into RAM from disk
  - Basic implementation: MMU with single-level page map
    - Access loaded pages via fast hardware path
    - Load virtual memory on demand: page faults
  - Several optimizations:
    - Moving page map to RAM, for cost reasons
    - Translation Lookaside Buffer (TLB) to regain performance
  - Cache/VM interactions: Can cache physical or virtual locations
- Goals 2 & 3: Ease memory management, protect multiple contexts from each other
  - We'll see these in detail on the next lecture!