


SVCs, Devices & Interrupts



- Interrupts for Async I/O
- ReadKey SVC example
- “Real Time”
 - Latencies & Deadlines
- Weak priorities
- Strong priorities

Today: Lab 6 due

Next week:

Wed: no recitation

Returning to User-mode

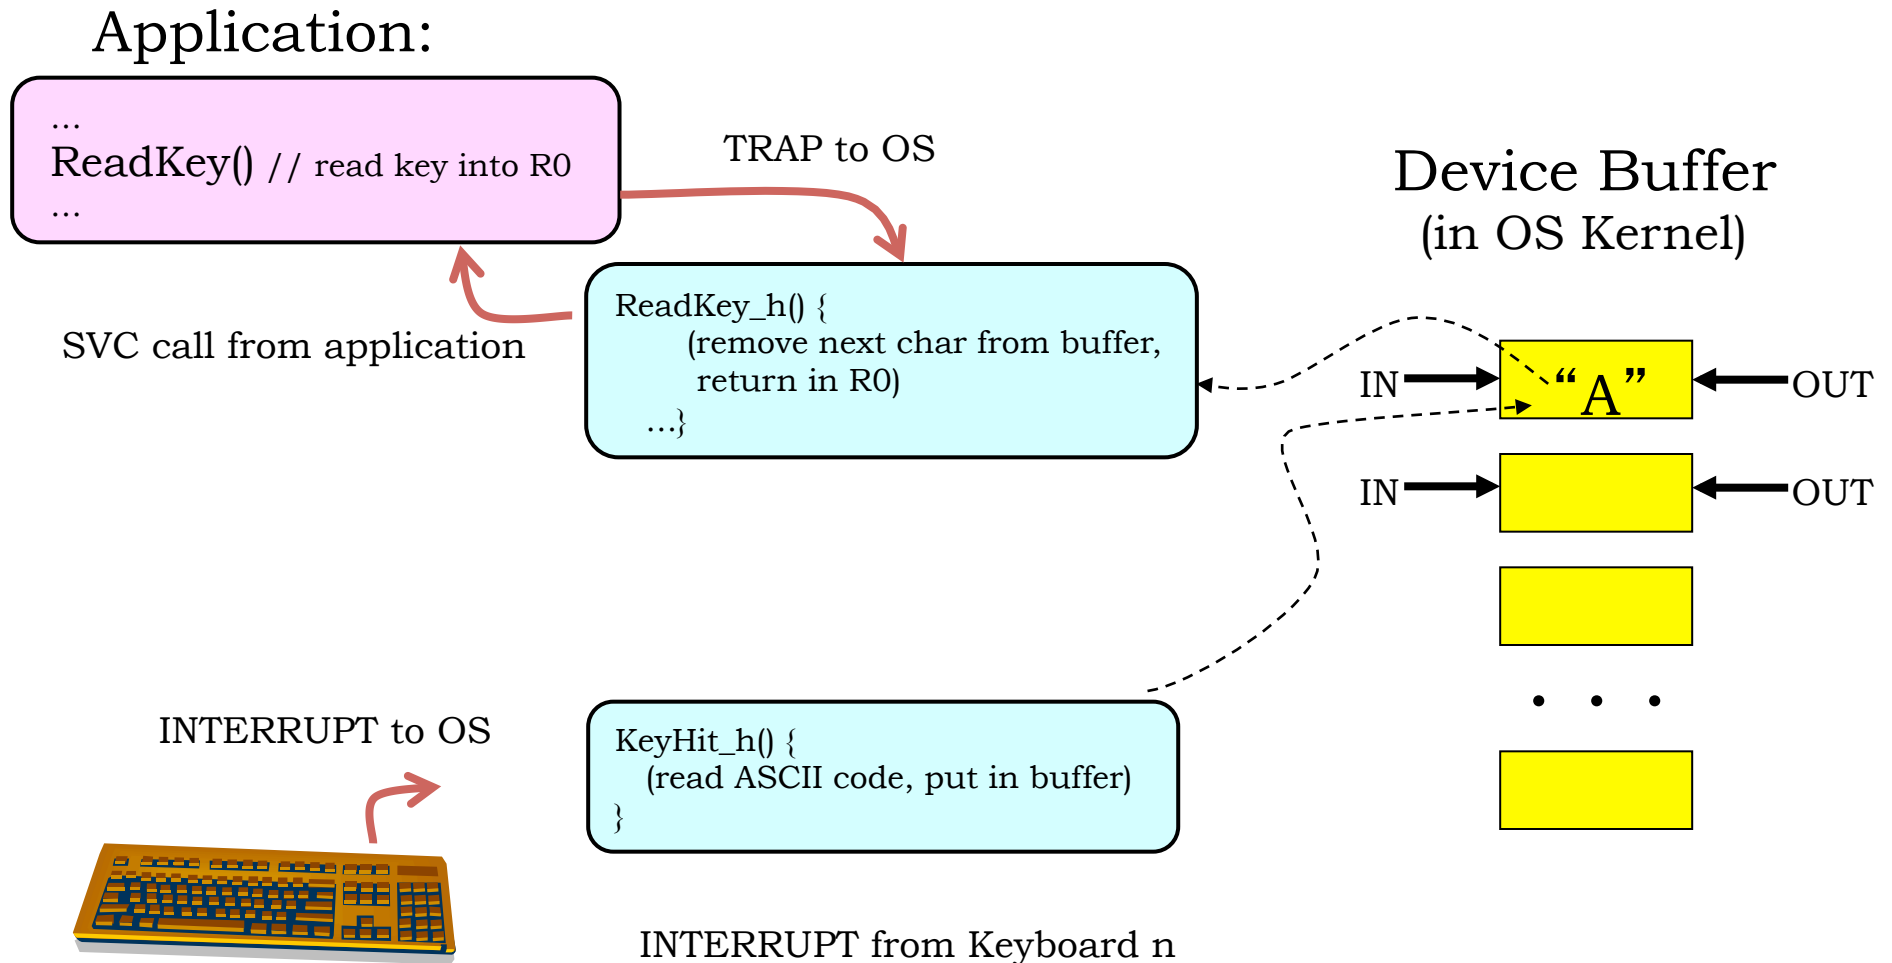
```
// Alternate return from interrupt handler which BACKS UP PC,  
// and calls the scheduler prior to returning. This causes  
// the trapped SVC to be re-executed when the process is  
// eventually rescheduled...
```

```
I_Wait: LD(UserMState+(4*30), r0) // Grab XP from saved MState,  
        SUBC(r0, 4, r0)           // back it up to point to  
        ST(r0, UserMState+(4*30)) // SVC instruction  
  
        CALL(Scheduler)          // Switch current process,  
                                  // and return to (some) user.
```

```
// Here's the common exit sequence from Kernel interrupt handlers:  
// Restore registers, and jump back to the interrupted user-mode  
// program.
```

```
I_Rtn:  RESTORESTATE()  
kexit:  JMP(XP)           // Good place for debugging breakpoint!
```

Asynchronous I/O Handling



Interrupt-based Asynch I/O


OPERATION: NO attention to Keyboard during normal operation

- on key strike: hardware asserts IRQ to request interrupt
- USER program interrupted, PC+4 of interrupted inst. saved in XP
- state of USER program saved on KERNEL stack;
- Keyboard handler invoked, runs to completion;
- state of USER program restored; program resumes.

TRANSPARENT to USER program.

Keyboard Interrupt Handler (in O.S. KERNEL):

Assume each keyboard has an associated buffer



```
struct Device {  
    char Flag, Data;  
} Keyboard;
```


```
KeyHit_h() {  
    Buffer[inptr] = Keyboard.Data;  
    inptr = (inptr + 1) % BUFSIZE;  
}
```

ReadKey SVC: Attempt #1

A *supervisor call* (SVC) is an instruction that transfers control to the kernel so it can satisfy some user request. Kernel returns to user program when request is complete.

First draft of a ReadKey SVC handler (supporting a *Virtual Keyboard*): returns next keystroke on a user's keyboard to that user's requesting application:

```
ReadKey_h()
{
    int kbdnum = ProcTbl[Cur].DPYNum;
    while (BufferEmpty(kbdnum)) {
        /* busy wait loop */
    }
    User.Reg[0] = ReadInputBuffer(kbdnum);
}
```



Problem: Can't interrupt code running in the supervisor mode... so the buffer never gets filled.

ReadKey SVC: Attempt #2

A BETTER keyboard SVC handler:

```
ReadKey_h()
{
    int kbdnum = ProcTbl[Cur].DPYNum;
    if (BufferEmpty(kbdnum)) {
        /* busy wait loop */
        User.Reg[XP] = User.Reg[XP]-4;
    } else
        User.Reg[0] = ReadInputBuffer(kbdnum);
}
```

That's a
funny way
to write
a loop



This one actually works!

Problem: The process just wastes its time-slice waiting for someone to hit a key...

ReadKey SVC: Attempt #3

EVEN BETTER: On I/O wait, YIELD remainder of quantum:

```
ReadKey_h()
{
    int kbdnum = ProcTbl[Cur].DPYNum;
    if (BufferEmpty(kbdnum)) {
        User.Reg[XP] = User.Reg[XP]-4;
        Scheduler( );
    } else
        User.Reg[0] = ReadInputBuffer(kbdnum);
}
```

RESULT: Better CPU utilization!!

Does timesharing cause CPU use to be less efficient?

- COST: Scheduling, context-switching overhead; but
- GAIN: Productive use of idle time of one process by running another.

Sophisticated Scheduling

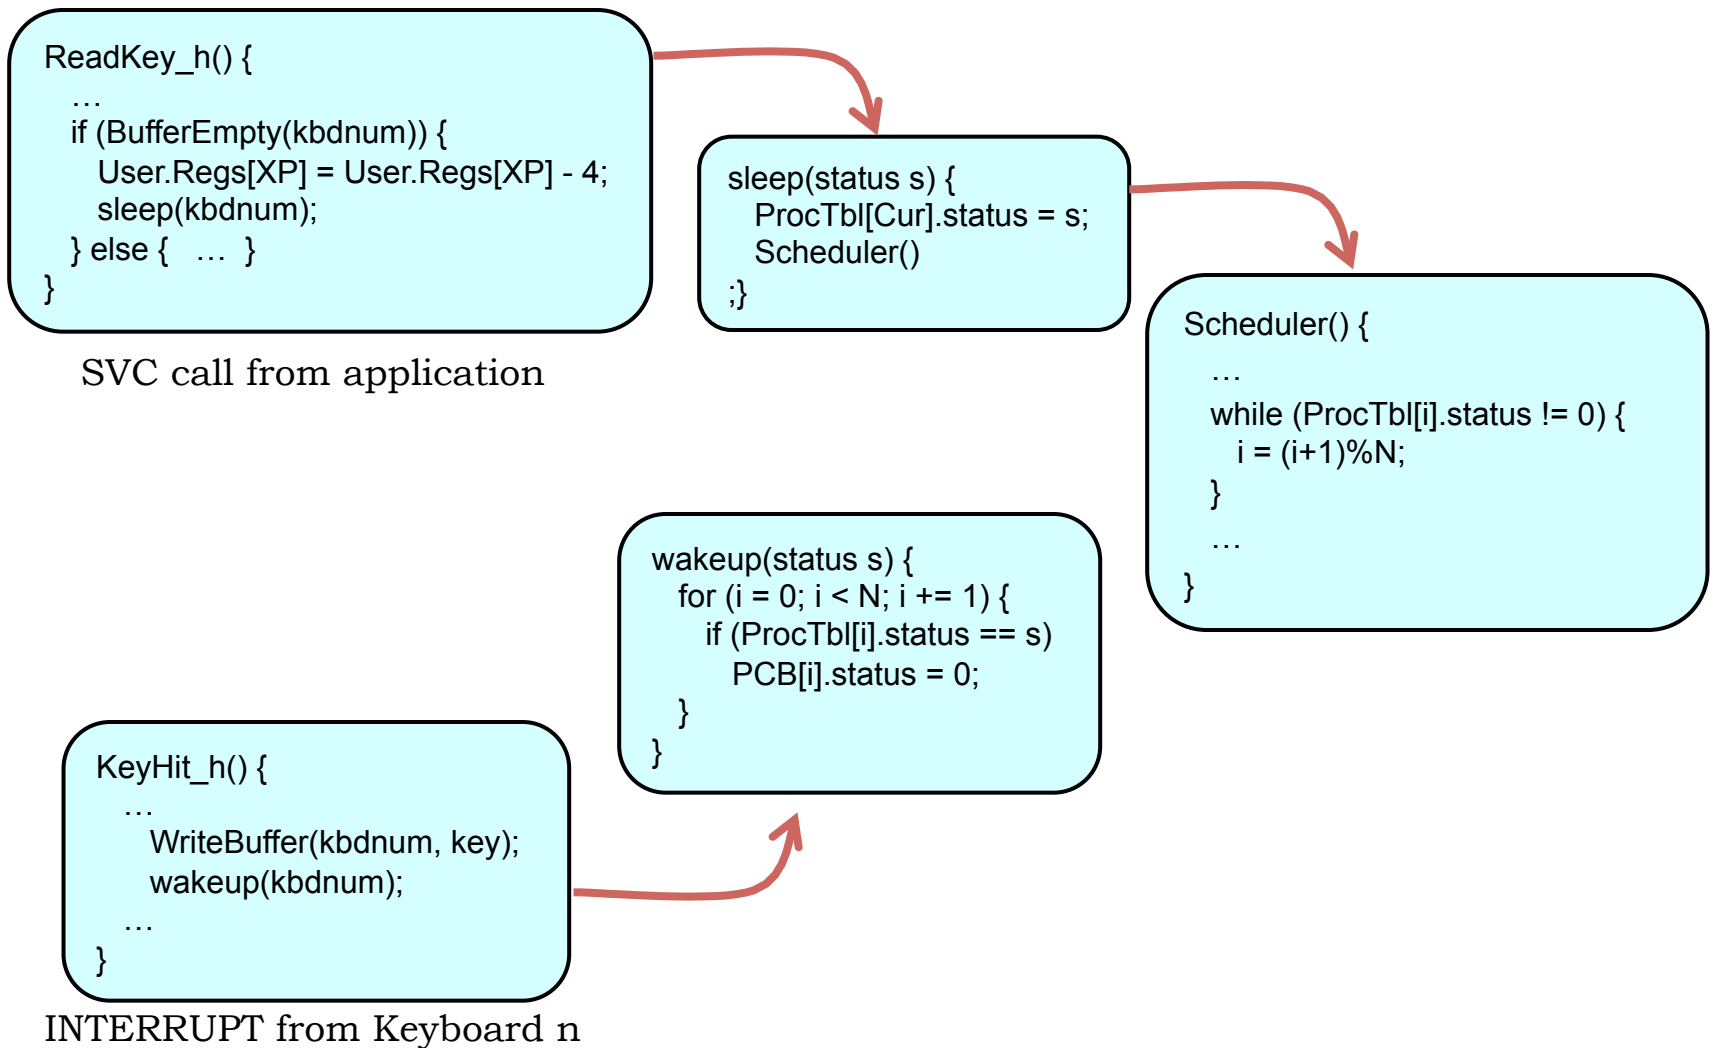
To improve efficiency further, we can avoid scheduling processes in prolonged I/O wait:

- Processes can be in **ACTIVE** or **WAITING** (“sleeping”) states;
- Scheduler cycles among **ACTIVE PROCESSES** only;
- Active process moves to **WAITING** status when it tries to read a character and buffer is empty;
- Waiting processes each contain a code (eg, in PCB) designating what they are waiting for (eg, keyboard N);
- Device interrupts (eg, on keyboard N) move any processes waiting on that device to **ACTIVE** state.

UNIX kernel utilities:

- `sleep(reason)` - Puts CurProc to sleep. “Reason” is an arbitrary binary value giving a condition for reactivation.
- `wakeup(reason)` - Makes active any process in `sleep(reason)`.

ReadKey SVC: Attempt #4



The Need for “Real Time”


Side-effects of CPU virtualization

- + abstraction of machine resources (memory, I/O, registers, etc.)
- + multiple “processes” executing concurrently
- + better CPU utilization
- Processing throughput is more variable




Our approach to dealing with the asynchronous world

- I/O - separate “event handling” from “event processing”



Difficult to meet “hard deadlines”

- control applications
- playing videos/MP3s




Real-time as an alternative to time-sliced or fixed-priority preemptive scheduling

Interrupt Latency

One way to measure the real-time performance of a system is *INTERRUPT LATENCY*:

- *HOW MUCH TIME can elapse between an interrupt request and the START of its handler?*




OFTEN bad things happen when service is delayed beyond some **deadline** - "real time" considerations:

Missed characters
System crashes
Nuclear meltdowns

} "HARD"
Real time
constraints



Sources of Interrupt Latency



What causes interrupt latency:

- State save, context switch.
- Periods of uninterruptability:

We can consider this when we write our O/S

We can address this in our ISA

But, this is application dependent!

- Long, uninterruptable instructions -- eg block moves, multi-level indirection.
- Explicitly disabled periods (eg for atomicity, during service of other interrupts).


GOAL: BOUND (and minimize) interrupt latency!

- Optimize interrupt sequence context switch
- Make unbounded-time instructions *interruptable* (state in registers, etc).
- Avoid/minimize disable time
- Allow handlers to be interrupted, in certain cases (while still avoiding **reentrant** handlers!).

Scheduling of Multiple Devices

"TOY" System scenario:	Actual w/c <u>Latency</u>	<u>DEVICE</u>	<u>Service Time</u>
	$500 + 400 = \underline{900}$	Keyboard	800
	$800 + 400 = \underline{1200}$	Disk	500
	$800 + 500 = \underline{1300}$	Printer	400

What is the WORST CASE latency seen by each device?




Assumptions:

- Infrequent interrupt requests (each happens only once/scenario)
 - Simultaneous requests might be served in ANY order.... Whence
 - Service of EACH device might be delayed by ALL others!
- ... can we improve this?

Weak (Non-preemptive) Priorities

ISSUE: Processor becomes interruptable (at fetch of next instruction), several interrupt requests are pending. Which is served first?



WEAK PRIORITY ORDERING: Check in prescribed sequence, eg: DISK > PRINTER > KEYBOARD.

Latencies with WEAK PRIORITIES:

Service of each device might be delayed by:

- Service of 1 other (arbitrary) device, whose interrupt request was just honored;
- +
- Service of ALL higher-priority devices.

Actual w/c Latency	DEVICE	Service Time
900	Keyboard	800
800	Disk	500
1300	Printer	400

vs 1200 –
Now delayed by only 1 service!

The Need for Preemption

Without preemption, ANY interrupt service can delay ANY other service request... the slowest service time constrains response to fastest devices. Often, tight deadlines can't be met using this scheme alone.

EXAMPLE: 800 uSec deadline (hence 300 uSec maximum interrupt latency) on disk service, to avoid missing next sector...

<u>Priority</u>	<u>Latency w/ preemption</u>	<u>Actual Latency</u>	<u>DEVICE</u>	<u>Serv. Time</u>	<u>Max. Delay</u>
1	D,P 900	900	Keybrd	800	300
3	~0	800	Disk	500	
2	[D] 500	1300	Printer	400	

CAN'T SATISFY the disk requirement in this system using weak priorities!

need **PREEMPTION**: Allow handlers for LOWER PRIORITY interrupts to be interrupted by HIGHER priority requests!

Strong Priority Implementation

STRONG PRIORITY ORDERING: Allow handlers for LOWER PRIORITY interrupts to be preempted (interrupted) by HIGHER PRIORITY requests.

SCHEME:

- Expand supervisor bit in PC to be a PRIORITY integer PRI (eg, 3 bits for 8 levels)
- ASSIGN a priority to each device.
- Prior to each instruction execution:
 - Find priority P_i of highest requesting device, say D_i
 - Take interrupt if and only if $P_i > \text{PRI}$, set $\text{PRI} = P_i$.

PC:

PRI	Program Counter
-----	-----------------

Strong priorities:

KEY: Priority in Processor state

Allows interruption of (certain) handlers

Allows preemption, but not reentrance

BENEFIT: Latency seen at high priorities UNAFFECTED by service times at low priorities.


Recurring Interrupts

Consider interrupts which recur at bounded rates:

<u>Actual Latency</u>	<u>DEVICE</u>	<u>P</u>	<u>Serv. Time</u>	<u>Max. Delay</u>	<u>Max. Freq</u>
900	Keybrd	1	800		100/s
0	Disk	3	500	300	500/s
500	Printer	2	400		1000/s

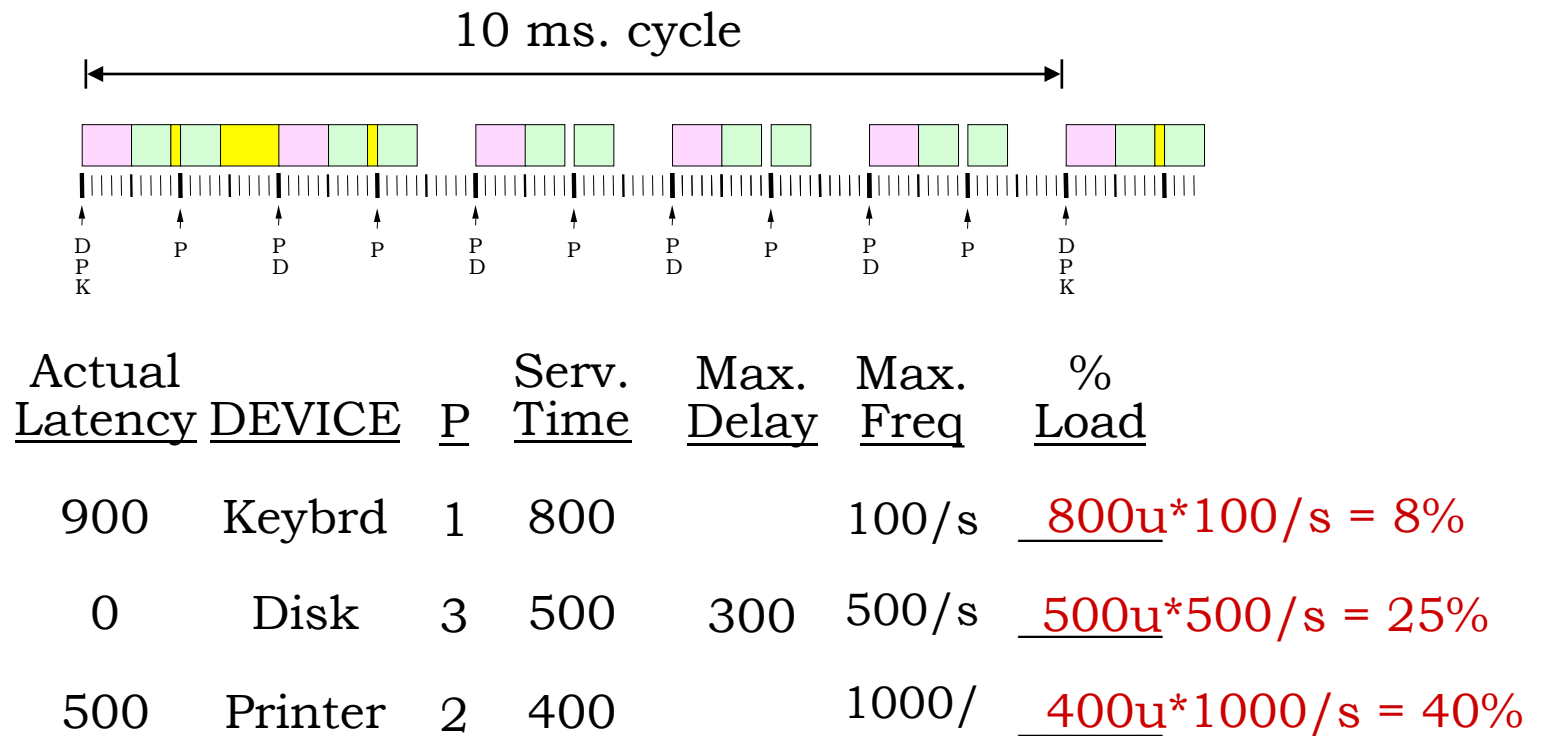
Note that interrupt LATENCIES don't tell the whole story—consider COMPLETION TIMES, eg for Keyboard in example to the right.

Keyboard service not complete until 3ms after request. Often *deadlines* used rather than max. delays.



Interrupt Load

How much CPU time is consumed by interrupt service?




Remaining fraction (27%) is left over for application; *trouble if it's <0!*

Example: Ben Visits the ISS



International Space Station's on-board computer performs 3 tasks:

- guiding incoming supply ships to a safe docking
- monitoring gyros to keep solar panels properly oriented
- controlling air pressure in the crew cabin



	<i>Task</i>	<i>Period</i>	<i>Service time</i>	<i>Deadline</i>	
16.6 %	Supply ship guidance	30ms	5ms	25ms	C,G = 10 + 10 + (5) = 25
25%	Gyroscopes	40	10	20	C = 10 + (10) = 20
10%	Cabin pressure	100	? 10	100	S,G = 5 + 10 + (10) = 25

Assuming a **weak priority system**:

1. What is the maximum service time for “cabin pressure” that still allows all constraints to be met? **< 10 mS**
2. Give a weak priority ordering that meets the constraints **G > SSG > CP**
3. What fraction of the time will the processor spend idle? **48.33%**
4. What is the worst-case delay for each type of interrupt until *completion* of the corresponding service routine?

Example: Ben Visits ISS (cont'd.)

Our Russian collaborators don't like the sound of a “weak” priority interrupt system and lobby heavily to use a “strong” priority interrupt system instead.

	<i>Task</i>	<i>Period</i>	<i>Service time</i>	<i>Deadline</i>	
16.6%	Supply ship guidance	30ms	5ms	25ms	[G] 10 + 5
25%	Gyroscopes	40	10	20	10
50%	Cabin pressure	100	? 50	100	100

Assuming a **strong priority system**, $G > \text{SSG} > \text{CP}$:

1. What is the maximum service time for “cabin pressure” that still allows all constraints to be met? $100 - (3 \cdot 10) - (4 \cdot 5) = 50$
2. What fraction of the time will the processor spend idle? **8.33%**
3. What is the worst-case delay for each type of interrupt until *completion* of the corresponding service routine?

Summary

Device interface – two parts:

- Device side: handle interrupts from device (transparent to apps)
- Application side: handle interrupts (SVCs) from application

Scheduler interaction:

- “Sleeping” (*inactive) processes waiting for device I/O
- Handler coding issues, looping thru User mode

Real Time constraints, scheduling, guarantees”

- Complex, hard scheduling problems – a black art!
- Weak (non-preemptive) vs Strong (preemptive) priorities help...
- Common real-world interrupt systems:
 - Fixed number (eg, 8 or 16) of strong priority levels
 - Each strong priority level can support many devices, arranged in a weak priority chain